

SAI AKHIL TEKURI

[LinkedIn](#) | +1 (716) 907-9073 | st5050@nyu.edu | [GitHub](#) | [Portfolio](#)

EDUCATION

New York University, Tandon School of Engineering, Brooklyn, New York 09/2023 - 05/2025
Master of Science, Computer Science, GPA: 3.889

Indian Institute of Technology Bombay, Mumbai, India 08/2018 - 05/2022
Bachelor of Technology with Honors, Computer Science and Engineering

SKILLS

- C++ | Python | Java | Scala | JavaScript | SQL | NodeJS | React | PostgreSQL | MongoDB | ClickHouse | Elasticsearch
- PyTorch | TensorFlow | LangGraph | LangChain | LLM | RAG | ChromaDB | Scikit-learn | OpenCV | PySpark | Kafka
- ETL | Data warehouse | AWS Microservices | Azure | Kubernetes | Docker | GenAI | Backend Development | Jira | Git

EXPERIENCE

Software Engineer, AI | *Gruve Tech Inc, Texas, U.S.A* 01/2026 - Present

- Architected an **AI FinOps** platform for **Stanford Health Care** on **Azure**, enabling enterprise-wide AI cost attribution
- Deployed **LiteLLM** with OpenTelemetry as an **LLM gateway** and built an **App Insights** to **Databricks** ETL pipeline with custom executive **FinOps dashboard** for each AI workflow to **track cost, model usage, and optimize AI spend**
- Architected an end-to-end **Insurance AI Platform**, automating claims processing and reducing manual effort by **50%**
- Developed a **LangGraph-based AI agent** integrating real-time **DuckDuckGo** search for insurance claim automation, covering damage assessment, cost estimation, fraud detection, and decision, **reducing claim processing time by 60%**
- Created a **Policy Agent** that extracts **adjudication** rules from client policies, enabling **customized AI claim decisions**
- Designed a **model-agnostic inference layer** enabling runtime switching among **OpenAI-compatible VLM** providers

Gen AI Engineer | *Infinite Computer Solutions, Texas, U.S.A* 06/2025 - 01/2026

- Designed **agentic AI system** for **Verizon** to automate the **SDLC workflow**, improving feature delivery speed by **40%**
- Created an **autonomous agent** with **LangGraph** to generate feature specifications from the business requirement doc with **70% accuracy**, leveraging a **Graph RAG pipeline** over a **Neo4j** enterprise knowledge graph with **20M+ nodes**
- Integrated a **conversational AI** assistant to support **system architects** and **dev teams** in refining the generated feature
- Achieved **context-aware, multi-session memory** assistant by integrating **LangGraph** checkpoints with **PostgreSQL**

Engineering Intern | *SecurityScorecard, New York, U.S.A* 01/2025 - 04/2025

- Delivered an **AI chatbot system** using **LangGraph** and **Claude 3.5 Sonnet** to **explain, contextualize** security scores
- Developed **RESTful APIs** in **Scala** to fetch score inputs, score outputs, and scoring context from **S3** and **ClickHouse**
- Integrated **custom tools** with **Claude** to improve **response accuracy** and to provide **clear insights** into security scores

AI Engineer Intern | *Radical AI, New York, U.S.A* 05/2024 - 08/2024

- Built an AI teaching assistant using **LangChain**, **Gemini 1.5 Pro**, and **ChromaDB** to generate concept-based quizzes
- Enhanced the AI's capability to support **15 new document types** by integrating **custom document loaders** in **Python**

Software Development Engineer | *Adobe Inc., Bangalore, India* 06/2022 - 07/2023

- Implemented **AWS-based** data ingestion by integrating 4 sources using the **ETL process** to build a **data warehouse**
- Revamped internal tools using **Scala**, **Docker**, and **Jenkins**, resulting in a **10% reduction** in **data processing time**
- Accomplished **20% reduction** in overall **data ingestion time** by **effective data transformation** and **storage format**
- Conducted in-depth analysis of **data marts and data vaults**, extracted **business insights**, and reduced cost by **10%**

PROJECTS

Multi-Tenant DBaaS with Automated Migrations | *Cloud Computing* 10/2024 - 12/2024

- **Architected** and **engineered** a scalable multi-tenant Database-as-a-Service framework with asynchronous migrations
- Leveraged **AWS Lambda, EC2, DynamoDB, SQS** to implement serverless workflows with **database-level isolation**

Predictive Modeling for Solar Energy Production in Europe | *Big Data* 02/2024 - 05/2024

- Trained a **Gradient Boosting Regressor** using **PySpark MLlib** to predict solar energy, attained a **RMSE of 0.236**
- Engineered real-time solar data streaming with **Kafka** for solar energy forecasting, visualized using **Tableau** dashboard